

## Lecture 1

*Instructor: Pedro Felzenszwalb**Scribes: Dan Xiang, Tyler Dae Devlin*

## Introduction

### Some logistical details

- Prerequisites: Linear algebra (MATH 0520), probability & statistics (APMA 1650), and introductory computer science (CSCI 0150, 0160).
- This course focuses on the mathematical foundations of machine learning. An alternative, more practically oriented course is CS1951A: Data Science.
- This course's website is <http://cs.brown.edu/~pff/engn2520/>.
- The class Piazza page is [www.piazza.com/brown/spring2017/cs142](http://www.piazza.com/brown/spring2017/cs142).

### Supervised learning

The goal of *supervised learning* is to estimate a function  $f : X \rightarrow Y$  where  $X$  is the “feature space” and  $Y$  is the “label space”. The function  $f$  is estimated using a set of labeled training data

$$T = \{(x_1, y_1), \dots, (x_n, y_n)\}$$

with  $\{x_i\} \subset X$  and  $\{y_i\} \subset Y$ , where each  $y_i = f(x_i)$  is the true label for the corresponding  $x_i$ .

Contrast this to *unsupervised learning*, in which examples don't come with a set of known labels, i.e. we don't specify an output space. Instead of trying to estimate a target function  $f$ , we just wish to analyze the structure of the data  $\{x_1, \dots, x_n\}$ .

*Example.* Let the input space be  $X = \mathbb{R}^2$ , let the output space be  $Y = \{\text{Bass}, \text{Salmon}\}$ , and suppose each example  $(x_1, x_2) \in X$  represents the length and weight of a fish. Given training data, we can estimate  $f : X \rightarrow Y$ , the function that determines whether a fish is a bass or salmon based on its length and weight. We can then use this estimate of  $f$  to classify new fish after weighing and measuring them.

Suppose we have some training data  $T$ , consisting of the lengths and weights of a bunch of fish along with their true species label. Define the classifier

$$g(x_1, x_2) = \begin{cases} \text{Salmon} & x_1 \geq t \\ \text{Bass} & x_1 < t \end{cases}.$$

The classifier  $g$  is an estimate of the target function  $f$ . Note that this is probably a poor estimate, since  $g$  classifies fish based only on their length, even though the training data also includes information about weight.

Nevertheless, to fully specify the classifier  $g$  we need to select the threshold  $t$ . Intuitively, we should select  $t$  to minimize the number of mistakes the classifier makes on the examples in the training data set, i.e. the **training error**. What we are really interested in, however, is the number of mistakes our function would make on a new set of unseen data, i.e. the **test error**. In general, it is easy to devise a classifier that has low (or even zero) training error; this is called *memorizing*. It is much harder to construct a classifier that generalizes to new data; this is called *learning*.  $\square$

**Theorem:** If the number of samples  $n$  is large enough, then the training error  $\approx$  test error with high probability. (We will make this precise later in the course.)

Leaving the fish classification example aside for the moment, let us examine a different kind of classifier. Suppose we have a training set  $T = \{(x_1, y_1), \dots, (x_n, y_n)\}$ . Define the classifier

$$g(x) = \begin{cases} y_j & x = x_j \text{ for some } j \in \{1, \dots, n\} \\ 0 & x \neq x_j \text{ for all } j \in \{1, \dots, n\} \end{cases}.$$

In words, this classifier checks if the feature  $x$  was observed in the training data. If it was, then it assigns the corresponding observed  $y_j$  label. Otherwise,  $g = 0$ .

*Example.* Classification is a very common task. A classic example is spam detection. In this case,

$$\begin{aligned} X &= \text{set of all possible emails,} \\ Y &= \{\text{spam, not spam}\}. \end{aligned}$$

We assume there is some true function  $f$  that associates to each email a label of either “spam” or “not spam”. Our goal is to use labeled training examples to construct an estimate  $g$  of  $f$  so that we can classify new emails and place them in the appropriate folders in your mail client.  $\square$

## Formalizing the problem

Given a feature space  $X$  and a label space  $Y$ , there is some distribution/density  $p(x, y)$  over  $X \times Y$  (the Cartesian product of  $X$  and  $Y$ ). We define a loss function  $L : Y \times Y \rightarrow \mathbb{R}$  by

$$L(y, \hat{y}) = \text{cost of predicting } \hat{y} \text{ if the true label is } y.$$

A quantity of interest is the expected loss of predictions using our learned classifier  $g$ , denoted

$$E[L(y, g(x))] \tag{1}$$

where  $y$  is the true label,  $g(x)$  is our predicted label, and  $(x, y) \sim p(x, y)$ .

We can define a particular loss function

$$L(y, \hat{y}) = \begin{cases} 0 & y = \hat{y} \\ 1 & y \neq \hat{y} \end{cases}.$$

Then equation (1) is the probability of error (since the expectation of an indicator of an event is precisely the probability of that event).

We typically assume that a training set  $T$  is a collection of iid samples drawn according to the joint distribution  $p(x, y)$ . Then we have that expression (1) is roughly equal to

$$\frac{1}{n} \sum_{i=1}^n L(y_i, g(x_i)),$$

i.e. the sample mean of  $L$ , also known as the empirical loss.

Given the underlying distribution  $p(x, y)$ , we compute using the definition of expected value

$$E[L(y, g(x))] = \int_X \sum_{y \in Y} p(x, y) L(y, g(x)) dx$$

For each  $x$ , we define the classifier

$$g(x) = \operatorname{argmin}_{\hat{y} \in Y} \sum_{y \in Y} p(x, y) L(y, \hat{y}).$$

In this case we can use the training data  $T$  to estimate the joint distribution  $p(x, y)$ . We can write

$$p(x, y) = p(y)p(x | y).$$

We can estimate  $p(y)$  by observing how many instances of each class  $y \in Y$  occur in  $T$ . For the  $p(x | y)$  term, we can estimate  $p(x | \text{Salmon})$  and  $p(x | \text{Bass})$  as Gaussians (for example) using parameters inferred from the training data.